

Brenda Estefania Castillo Fernandez

```
[ ] import pandas as pd
path = "/content/Social_Network_Ads.csv"
df = pd.read_csv(path)
```

```
[ ] df.head(5)
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0

```
[ ] df.shape
```

(400, 5)

```
[ ] df.describe()
```

	User ID	Age	EstimatedSalary	Purchased
count	4.000000e+02	400.000000	400.000000	400.000000
mean	1.569154e+07	37.655000	69742.500000	0.357500
std	7.165832e+04	10.482877	34096.960282	0.479864
min	1.556669e+07	18.000000	15000.000000	0.000000
25%	1.562676e+07	29.750000	43000.000000	0.000000
50%	1.569434e+07	37.000000	70000.000000	0.000000
75%	1.575036e+07	46.000000	88000.000000	1.000000
max	1.581524e+07	60.000000	150000.000000	1.000000

```
[ ] print(df)
```

	User ID	Gender	Age	EstimatedSalary	Purchased
0	15624510	Male	19	19000	0
1	15810944	Male	35	20000	0
2	15668575	Female	26	43000	0
3	15603246	Female	27	57000	0
4	15804002	Male	19	76000	0
...
395	15691863	Female	46	41000	1
396	15706071	Male	51	23000	1

```
[ ] df.drop('Gender', axis=1) #el drop se utiliza para quitar una columna. tmbn se puede hacer con columns='Gender'
```

	User ID	Age	EstimatedSalary	Purchased
0	15624510	19	19000	0
1	15810944	35	20000	0
2	15668575	26	43000	0
3	15603246	27	57000	0
4	15804002	19	76000	0
...
395	15691863	46	41000	1
396	15706071	51	23000	1
397	15654296	50	20000	1
398	15755018	36	33000	0
399	15594041	49	36000	1

400 rows x 4 columns

se necesita una label, la cual, en este caso es el purchased, la label es yes la y por quienes ya compraron, quienes no han comprado por sus características. las demás columnas son x

Brenda Estefania Castillo Fernandez

```
[ ] dfxtrain = df.iloc[:319] #este es el 80% del dataset

[ ] print(dfxtrain.shape)

(319, 5)

[ ] dfxtest = df.iloc[320:] #este es el 20%

[ ] dfxtest.shape

(80, 5)

[ ] y = df['Purchased']

[ ] print(y)

0      0
1      0
2      0
3      0
4      0
..
395    1
396    1
397    1
398    0
399    1
Name: Purchased, Length: 400, dtype: int64

[ ] ytrain = y.iloc[:319]

[ ] ytest = y.iloc[320:]

[ ] print(ytrain.shape)

(319,)

[ ] print(ytest.shape)

(80,)

[ ] import sklearn

[ ] df['Male'] = (df['Gender'] == 'Male').astype(int)
df['Female'] = (df['Gender'] == 'Female').astype(int)
df = df.drop(columns=['Gender'])

[ ] df = df[['User ID', 'Male', 'Female', 'Age', 'EstimatedSalary', 'Purchased']]

[ ] df.drop('User ID', axis=1, inplace=True) #esto es para eliminar la columna del user id del data set

[ ] df.head(15)
```

Brenda Estefania Castillo Fernandez

```
[ ]
```

	Male	Female	Age	EstimatedSalary	Purchased
0	1	0	19	19000	0
1	1	0	35	20000	0
2	0	1	26	43000	0
3	0	1	27	57000	0
4	1	0	19	76000	0
5	1	0	27	58000	0
6	0	1	27	84000	0
7	0	1	32	150000	1
8	1	0	25	33000	0
9	0	1	35	65000	0
10	0	1	26	80000	0
11	0	1	26	52000	0
12	1	0	20	86000	0
13	1	0	32	18000	0
14	1	0	18	82000	0

```
[ ] dfxtrain = df.iloc[:319, 0:4]
```

```
[ ] dfxtest = df.iloc[320:, 0:4]
```

```
[ ] y = df["Purchased"]
```

```
print(dfxtest)
print(dfxtrain)
```

```
Male  Female  Age  EstimatedSalary
320    0      1  52         138000
321    0      1  53          82000
322    1      0  41          52000
323    0      1  48          30000
324    0      1  48         131000
..     ...   ...   ...
395    0      1  46          41000
396    1      0  51          23000
397    0      1  50          20000
398    1      0  36          33000
399    0      1  49          36000
```

```
[80 rows x 4 columns]
```

```
Male  Female  Age  EstimatedSalary
0     1      0  19          19000
1     1      0  35          20000
2     0      1  26          43000
3     0      1  27          57000
4     1      0  19          76000
..     ...   ...   ...
314    0      1  39          79000
315    0      1  39          75000
316    0      1  54         104000
```

Brenda Estefania Castillo Fernandez

```
[ ] from sklearn.linear_model import Perceptron
    #dfxtrain, dfy = load_digits(return_X_y=True)
    clf = Perceptron(tol=1e-3, random_state=0)
    clf.fit(dfxtrain,ytrain)
    clf.score(dfxtrain,ytrain) #es para tener el accuracy, la metrica de predicción, en este caso tiene el 71% de exactitud al separar un modelo de otro.
```

```
0.7115987460815048
```

```
[ ] clf.predict(dfxtest)
```

```
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
       0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```

```
[ ] print(ytest)
```

```
320 1
321 1
322 0
323 1
324 1
..
395 1
396 1
397 1
398 0
399 1
Name: Purchased, Length: 80, dtype: int64
```